

- Week 3: Project Walkthru: Connecting to Redshift
- Data Ingestion Practices
- Introducing to Google BigQuery

AWS: Launching a Redshift Cluster

วิธีใช้งาน AWS Management Console ผ่าน [AWS Learner Lab](#)

<https://youtu.be/T8noPec75po>

วิธีสร้าง Amazon Redshift cluster ขึ้นมาใช้งาน

<https://youtu.be/BXBb3E95I6Y>

วิธีการทำ Data Ingestion จาก S3 เข้า Amazon Redshift

ข้อมูล Titanic

<https://github.com/datasciencedojo/datasets/blob/master/titanic.csv>

ตัวอย่างโค้ดที่ใช้ในการสร้างตารางใน Redshift

```
CREATE TABLE IF NOT EXISTS titanic (
  PassengerId text primary key,
  Survived text,
  Pclass text,
  Name text,
  Sex text,
  Age text,
  SibSp text,
  Parch text,
  Ticket text,
  Fare text,
  Cabin text,
  Embarked text
)
```

ตัวอย่างโค้ดที่ใช้ในการโหลดข้อมูลจาก S3 เข้า Redshift

```
COPY titanic FROM 's3://zkan-swu-labs/titanic.csv'
CREDENTIALS 'aws_iam_role=arn:aws:iam::377290081649:role/LabRole'
IGNOREHEADER 1 CSV REGION 'us-east-1'
```

ตัวอย่างโค้ดในการทำ Transformation โดยใช้ SQL

```
INSERT INTO
  new_table (
    name
  )
```

```
SELECT
  Name
FROM
  titanic
WHERE
  Name NOT IN (SELECT DISTINCT name FROM new_table)
```

Last modified: Saturday, 10 September 2022, 2:21 PM

GCP: Setting up BigQuery

Google Cloud Platform: <https://cloud.google.com/>

ข้อมูล Titanic

<https://github.com/datasciencedojo/datasets/blob/master/titanic.csv>

ตัวอย่างโค้ดที่ใช้ในการสร้างตารางใน BigQuery

```
CREATE TABLE IF NOT EXISTS `swu.titanic` (
  PassengerId STRING,
  Survivedt STRING,
  Pclass STRING,
  Name STRING,
  Sex STRING,
  Age STRING,
  SibSp STRING,
  Parch STRING,
  Ticket STRING,
  Fare STRING,
  Cabin STRING,
  Embarked STRING
)
```

ตัวอย่างโค้ดที่ใช้ในการโหลดเข้า BigQuery จาก Google Cloud Storage (GCS)

```
LOAD DATA OVERWRITE `swu.titanic`
FROM FILES (
  format = 'CSV',
  uris = ['gs://bucket/path/file.csv']
);
```

ตัวอย่างโค้ดในการทำ Transformation โดยใช้ SQL

```
MERGE swu_staging.titanic as dest
USING swu.titanic as src
ON dest.PassengerId = src.PassengerId
WHEN NOT MATCHED THEN
  INSERT (PassengerId, Name)
  VALUES (src.PassengerId, src.Name)
```

WHEN MATCHED THEN

UPDATE SET

dest.PassengerId = src.PassengerId

, dest.Name = src.Name

ดูรายละเอียดเพิ่มเติมเกี่ยวกับ Merge statement ใน BigQuery ได้ที่ [Data manipulation language \(DML\) statements in Google Standard SQL](#)

Last modified: Saturday, 10 September 2022, 2:33 PM

Week 3

<https://youtu.be/cydYqx5UMLo>

Week 3 Office Hour

<https://youtu.be/DXCVY8obmZO>

git config pull.rebase false

เสร็จแล้วก็ git pull

create table if not exists category_staging (

catname text

);

insert into category_staging (catname)

select

case

when catname = 'NFL' then 'Hello World'

end

from category

where catname not in (select distinct catname from category_staging);

select * from category_staging

delete from category_staging where catname = 'NFL'

User: introde413@gmail.com

Pass: 5qGn@0udHZ!7XW9U7fh8Sp

Week 4 (Part 1)

https://youtu.be/iVlkcoe_IGc

Week 4 (Part 2)

<https://youtu.be/BKFPZkwSUKc>

Project: Building an ETL Pipeline for a Cloud Data Warehouse (Redshift/Google BigQuery)

สามารถเริ่มต้นจากโค้ดที่ <https://github.com/zkan/swu-ds525> ได้

ในโปรเจกต์นี้เราจะใช้ข้อมูล GitHub event data จาก [API](#) หรือสามารถดาวน์โหลดไฟล์ JSON ได้ที่ [URL](#) นี้
สิ่งที่คาดหวังในโปรเจกต์นี้

ในโปรเจกต์นี้เราจะตั้งต้นว่าเรามีไฟล์อยู่บน S3 อยู่แล้ว ดังนั้นทุกคนสามารถที่จะอัปโหลดไฟล์ขึ้นไปไว้บน S3 ได้เลย

1. มีโค้ดที่ทำ ETL จากข้อมูล JSON โหลดเข้าไปใน Redshift ตาม data model ที่เราออกแบบไว้
2. มีการเขียน documentation อธิบายสิ่งที่ตัวเองทำลงไป รวมไปถึงการออกแบบ data model
3. มี instruction ในการรันโค้ดของตัวเอง


อ่านวิธีโหลด JSON เข้า Redshift ได้ที่ [COPY from JSON format](#)

Last modified: Saturday, 3 September 2022, 8:05 AM

```
COPY titanic FROM 's3://zkan-swu-labs/titanic.csv'  
CREDENTIALS 'aws_iam_role=arn:aws:iam::377290081649:role/LabRole'  
IGNOREHEADER 1 CSV REGION 'us-east-1'
```

Data Ingestion Practices #1

(Conceptually)



ข้อดี 🍌

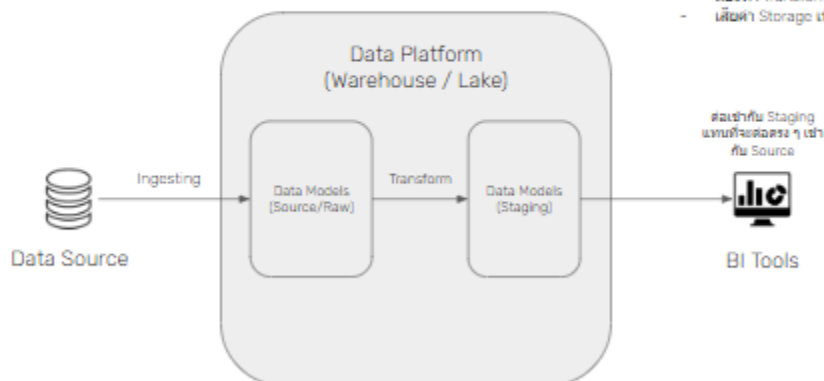
- ง่าย
- สะดวก
- ใช้ BI จัดการได้หมด

ข้อเสีย 🍌

- ข้อมูลไม่ได้คุณภาพ
- ข้อมูลผิด Privacy
- จัดการสิทธิ์การเข้าถึงข้อมูลได้ยาก

Data Ingestion Practices #2 (Better!)

(Conceptually)



ข้อดี 🍌

- ข้อมูลมีคุณภาพ
- สามารถจัดการข้อมูล Privacy ออกได้
- จัดการสิทธิ์การเข้าถึงข้อมูลได้

ข้อเสีย 🍌

- ต้องทำ Transform เพิ่ม
- เสียค่า Storage เพิ่ม

← → ↻

awsacademy.instructure.com/courses/24728/modules/items/2033633

🔖 ⭐ ⚙ ⌵ 👤 ⋮

🔍 poll_form.php

🔍 Canny edge detecti...

🔍 Simple guide on ho...

🔍 A Neural Network P...

🔍 Learner Lab

🔍 Image Feature Extra...

🔍 Canny edge detecti...

🔍 How to tune hyper...

🔍 Machine Learning u...

»

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

⏪

☰ ALLv1-24728 > Modules > Learner Lab > Learner Lab

Home

Modules

Discussions

AWS ●

Used \$3.3 of \$100

03:59

▶ Start Lab

■ End Lab

ℹ AWS Details

ℹ Readme

↺ Reset

✕

ddd_v1_w_foQ_1383611@runweb63672:-\$ []

EN-US -

Learner Lab

[Environment Overview](#)

[Environment Navigation](#)

[Access the AWS Management Console](#)

[Region restriction](#)

[Service usage and other restrictions](#)

[Using the terminal in the browser](#)

[Running AWS CLI commands](#)

[Using the AWS SDK for Python](#)

[Preserving your budget](#)

[Accessing EC2 Instances](#)

[SSH Access to EC2 Instances](#)

[SSH Access from Windows](#)

[SSH Access from a Mac](#)

Environment Overview

◀ Previous

Next ▶

us-east-1.console.aws.amazon.com/console/home?region=us-east-1#

poll_form.phpCanny edge detecti...Simple guide on ho...A Neural Network P...Learner LabImage Feature Extra...Canny edge detecti...How to tune hyper...Machine Learning u...

awsServicesSearch for services, features, blogs, docs, and more[Alt+S]N. Virginiavoclabs/user1591029=peeyapak.somvittoon@g.swu.ac.th @ 0748-312...

Console HomeInfo

Reset to default layoutAdd widgets

Recently visitedInfo

S3

EC2

EMR

Amazon Redshift

View all services

Welcome to AWS

Getting started with AWS

Learn the fundamentals and find valuable information to get the most out of AWS.

Training and certification

Learn from AWS experts and advance your skills and knowledge.

What's new with AWS?

Discover new AWS services, features, and Regions.

AWS HealthInfo

Open issues

0

Past 7 days

Scheduled changes

0

Upcoming and past 7 days

Other notifications

0

Past 7 days

Go to AWS Health

Cost and usageInfo

Current month costs

\$0.00

Forecasted month end costs

\$0.00

Down 100% from last month

Last month costs

\$3.30

Costs shown are unblended. [Learn more](#)

Top costs for current month

No top cost breakdowns. Cost breakdowns show when you use services.

Go to AWS Cost Management

Feedback

Looking for language selection? Find it in the new [Unified Settings](#)

© 2022, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

←

→

↻

us-east-1.console.aws.amazon.com/console/home?region=us-east-1#

poll_form.php

Canny edge detecti...

Simple guide on ho...

A Neural Network P...

Learner Lab

Image Feature Extra...

Canny edg

aws

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia ▾

Recently visited

Favorites

All services

Analytics

Application Integration

AR & VR

AWS Cost Management

Blockchain

Business Applications

Compute

Containers

Customer Enablement

Database

Developer Tools

End User Computing

Front-end Web & Mobile

Game Development

Internet of Things

Machine Learning

Management & Governance

Media Services

Migration & Transfer

Networking & Content Delivery

Quantum Technologies

Recently visited

Console Home

S3

Scalable Storage in the Cloud

EC2

Virtual Servers in the Cloud

EMR

Managed Hadoop Framework

Amazon Redshift

Fast, Simple, Cost-Effective Data Warehousing

[t+S]

N. Virginia ▲

US East (N. Virginia)

us-east-1

Recently visited

Favorites

All services

Analytics

Application Integration

AR & VR

AWS Cost Management

Blockchain

Business Applications

Compute

Containers

Customer Enablement

Database

Developer Tools

End User Computing

Front-end Web & Mobile

Game Development

Internet of Things

Machine Learning

 Management &
Governance

Media Services

Migration & Transfer

 Networking & Content
Delivery

AWS Data Exchange

Easily exchange data in the cloud

Data Pipeline

Orchestration for Data-Driven Workflows

EMR

Managed Hadoop Framework

Amazon FinSpace

Store, catalog, prepare, and analyze financial industry data

AWS Glue

AWS Glue is a fully managed ETL (extract, transform, and load) service

AWS Glue DataBrew

Visual data preparation tool to clean and normalize data for analytics and machine learning

Kinesis

Work with Real-Time Streaming Data

AWS Lake Formation

AWS Lake Formation makes it easy to set up a secure data lake

MSK

Fully managed, highly available, and secure service for Apache Kafka

Amazon OpenSearch Service

Run and Scale OpenSearch and Elasticsearch Clusters (successor to Amazon Elasticsearch Service)

QuickSight

Fast, easy to use business analytics

Amazon Redshift

Fast, Simple, Cost-Effective Data Warehousing



Analytics

Amazon Redshift

Accelerate your time to
insights with fast, easy, and
secure analytics at scale.

Amazon Redshift makes it easier for you to run and scale analytics without having to manage your data warehouse. Get insights by running real-time and predictive analytics on all of your data, across operational databases, data lake, data warehouse, and thousands of third-party datasets.

us-east-1.console.aws.amazon.com/redshiftv2/home?region=us-east-1#landing

poll_form.php Canny edge detecti... Simple guide on ho... A Neural Network P... Learner Lab Image Feature Extra... Canny edge detecti... How to tune hyper...

aws Services Search for services, features, blogs, docs, and more [Alt+S] N. Virginia voclabs/user1591029=peeyapak.somvitoo

Amazon Redshift

Accelerate your time to insights with fast, easy, and secure analytics at scale.

Amazon Redshift makes it easier for you to run and scale analytics without having to manage your data warehouse. Get insights by running real-time and predictive analytics on all of your data, across operational databases, data lake, data warehouse, and thousands of third-party datasets.


Get to powerful insights fast

The Amazon Redshift serverless experience makes it easy for customers to run and scale analytics without having to provision and manage their data warehouse. Simply load and query data.

[Try Amazon Redshift Serverless](#)

How it works

Introduction to Data Warehousing on AWS with ... [Copy link](#)



Amazon Redshift

Watch on YouTube

Redshift powers mission critical analytical workloads for Fortune 500 companies.

Provision and manage clusters

With a few clicks, you can create your first Amazon Redshift provisioned cluster in minutes.

[Create cluster](#)

Pricing and cost

[On-demand pricing](#)

[Reserved instance pricing](#)

ra3.4xlarge

Managed storage: up to 128 TB/node
\$3.26/node/hour
\$0.024/GB/month
12 vCPU (gen 3)

Amazon Redshift > Clusters > Create cluster

C

RA3 (recommended)

High performance with scalable managed storage

ra3.xlplus

Managed storage: up to 4 TB for single-node clusters. For clusters with 2 or more nodes, each node has storage up to 32 TB.
\$1.086/node/hour
\$0.024/GB/month
4 vCPU (gen 3)

Cluster configuration

Cluster identifier

This is the unique key that identifies a cluster.

The identifier must be from 1-63 characters. Valid characters are a-z (lowercase only) and - (hyphen).

What are you planning to use this cluster for?



Production

Configure for fast and consistent performance at the best price.



Free trial

Configure for learning about Amazon Redshift. This configuration is free for a limited time if your organization has never created an Amazon Redshift cluster.

Choose the size of the cluster

Node type [Info](#)

Choose a node type that meets your CPU, RAM, storage capacity, and drive type requirements.

Number of nodes

Enter the number of nodes that you need.

Range (1-32)

Configuration summary [Info](#)

ra3.xlplus | 1 node

Sample data [Info](#)



Load sample data

Load sample data to your Redshift cluster to start using the query editor to query data.

Tickit (28 MB)

Tickit is the sample data set that uses a sample database called TICKIT. Tickit contains individual sample data files: two fact tables and five dimensions.

Associated IAM roles (0) [Info](#)

Create, associate, or remove an IAM role. You can associate up to 50 IAM roles. You can also choose an IAM role and set it as the default for this cluster.

Set default ▼

Manage IAM roles ▼

🔍 Search for associated IAM role by name, status, or role type

< 1 >



IAM roles [🔗](#)



Status



Role type



No resources

No associated IAM roles

Associate IAM role

Associate IAM roles



IAM roles (1/3)

Choose from existing IAM roles. You can associate up to 50 IAM roles with this cluster.

🔍 Search for IAM role to associate

< 1 >



IAM roles [🔗](#)



AWSServiceRoleForRedshift



LabRole



myRedshiftRole

Cancel

Associate IAM roles

Additional configurations ☒ Use defaults



These configurations are optional, and default settings have been defined to help you get started with your cluster. Turn off "Use defaults" to modify these settings now.

▼ Network and security [Info](#)

Virtual private cloud (VPC)

This VPC defines the virtual networking environment for this cluster.


Default VPC
vpc-0644f8397501301da

 You can't change the VPC associated with this cluster after the cluster has been created. [Learn more](#) 

VPC security groups

This VPC security group defines which subnets and IP ranges the cluster can use in the VPC.

Choose one or more security groups

default 
sg-0b7184600e1e31554

Cluster subnet group

Choose the Amazon Redshift subnet group to launch the cluster in.

cluster-subnet-group-1

Availability Zone

Specify the Availability Zone to create the cluster in. Otherwise, Amazon Redshift chooses an Availability Zone for you.

No preference

Enhanced VPC routing

Enabling this option routes network traffic between your cluster and data repositories through a VPC, instead of through the internet. [Learn more](#) 

- ☒ Turn off
☐ Turn on

Amazon Redshift

×

Redshift serverless New

Provisioned clusters dashboard

▶ Clusters

Query editor

Query editor v2

Queries and loads

Datashares

▼ Configurations

Workload management

Subnet groups

HSM

Manage Tags

Advisor

AWS Marketplace

Alarms

Events

What's new

Amazon Redshift > Configurations > Subnet groups

Cluster subnet groups (1)

Create cluster subnet group

< 1 >

<input type="checkbox"/>	Name ▲	Status ▼	VPC ID ▼	Description ▼	Tags ▼
<input type="checkbox"/>	cluster-subnet-group-1 6 Subnets	Complete	vpc-0644f8397501301da	My subnet group	

aws

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

voclabs/user15910

Amazon Redshift

Redshift serverless

Provisioned clusters dashboard

Clusters

Query editor

Query editor v2

Queries and loads

Datashares

Configurations

Workload management

Subnet groups

HSM

Manage Tags

Advisor

AWS Marketplace

Alarms

Events

What's new

Amazon Redshift > Configurations > Subnet groups > Create subnet group

Create cluster subnet group

Cluster subnet group details

Name

You can't modify the name after your subnet group has been created.

cluster-subnet-group-2

The name must be 1-255 characters. Valid characters are A-Z, a-z, 0-9, space, hyphen (-), underscore (_), and period (.).

Description

Add subnets

VPC

Choose the VPC that contains the subnets that you want to include in your cluster subnet group.

Choose a VPC

Add all the subnets for this VPC

Availability Zone

Subnet

Choose an Availability Zone

Choose a subnet

Add subnet

Subnets in this cluster subnet group (0)

Remove all

Availability Zone	Subnet ID	CIDR block	Action

Create cluster subnet group

Cluster subnet group details

Name

You can't modify the name after your subnet group has been created.

The name must be 1-255 characters. Valid characters are A-Z, a-z, 0-9, space, hyphen (-), underscore (_), and period (.).

Description

Add subnets

VPC

Choose the VPC that contains the subnets that you want to include in your cluster subnet group.

vpc-0644f8397501301da



Add all the subnets for this VPC

Availability Zone



Subnet



Add subnet

Create cluster subnet group

Cluster subnet group details

Name

You can't modify the name after your subnet group has been created.

cluster-subnet-group-2

The name must be 1-255 characters. Valid characters are A-Z, a-z, 0-9, space, hyphen (-), underscore (_), and period (.).

Description

My subnet group.

Add subnets

VPC

Choose the VPC that contains the subnets that you want to include in your cluster subnet group.

Default VPC

vpc-0644f8397501301da

Add all the subnets for this VPC

Availability Zone

Choose an Availability Zone

Subnet

Choose a subnet

Add subnet

Subnets in this cluster subnet group (6)

Remove all

Availability Zone	Subnet ID	CIDR block	Action
us-east-1e	subnet-0c91d47d5ada65755	172.31.48.0/20	Remove
us-east-1c	subnet-05ef302e895205a56	172.31.16.0/20	Remove
us-east-1f	subnet-0a22a17cd312ea56b	172.31.64.0/20	Remove
us-east-1a	subnet-08ff834d363777a16	172.31.0.0/20	Remove
us-east-1b	subnet-04ed0369ff48942ed	172.31.80.0/20	Remove
us-east-1d	subnet-0cb36ac9dd90a44f9	172.31.32.0/20	Remove

Cancel

Create cluster subnet group

Additional configurations ☒ Use defaults




These configurations are optional, and default settings have been defined to help you get started with your cluster. Turn off "Use defaults" to modify these settings now.

▼ Network and security [Info](#)

Virtual private cloud (VPC)

This VPC defines the virtual networking environment for this cluster.


Default VPC
vpc-0644f8397501301da

 You can't change the VPC associated with this cluster after the cluster has been created. [Learn more](#)  

VPC security groups

This VPC security group defines which subnets and IP ranges the cluster can use in the VPC.

Choose one or more security groups

default
sg-0b7184600e1e31554 

Cluster subnet group

Choose the Amazon Redshift subnet group to launch the cluster in.

cluster-subnet-group-2

Availability Zone

Specify the Availability Zone to create the cluster in. Otherwise, Amazon Redshift chooses an Availability Zone for you.

No preference

Enhanced VPC routing

Enabling this option routes network traffic between your cluster and data repositories through a VPC, instead of through the internet. [Learn more](#) 

- ☒ Turn off
☐ Turn on

▼ Database configurations

Database name

Specify a database name to create an additional database.

The name must be 1-64 alphanumeric characters (lowercase only), and it can't be a reserved word.

Database port

Port number where the database accepts inbound connections. You can't change the port after the cluster has been created.

The port must be numeric (1150-65535).

Parameter groups

Defines database parameter and query queues for all the databases.

Default parameter group for redshift-1.0

Encryption

Encrypt all data on your cluster.

- ☒ Disabled
- ☐ Use AWS Key Management Service (AWS KMS)
- ☐ Use a hardware security module (HSM)

← → ↺

us-east-1.console.aws.amazon.com/redshiftv2/home?region=us-east-1#clusters

🔍 📄 ⭐ ⚙️ 🗑️ 👤

poll_form.php Canny edge detecti... Simple guide on ho... A Neural Network P... Learner Lab Image Feature Extra... Canny edge detecti... How to tune hyper... Machine Learning u...

aws Services 🔍 Search for services, features, blogs, docs, and more [Alt+S] 📧 🔔 ⓘ N. Virginia voclabe/user1591029=peeyapak.somvittoon@g.swu.ac.th @ 0748-312...

Amazon Redshift ×

Redshift serverless **New**

Provisioned clusters dashboard

▼ Clusters

Reserved nodes

Snapshots

Query editor

Query editor v2 **🔗**

Queries and loads

Dashboards

▼ Configurations

Workload management

Subnet groups

HSM

Manage Tags

Advisor

AWS Marketplace

Alarms

Events

What's new

Cluster subnet group cluster-subnet-group-2 was create successfully

redshift-cluster-1 is being created.

Sample data load

After the cluster is created, Amazon Redshift starts to load the sample data.

Amazon Redshift > Clusters

In my account From other accounts

▼ Connect to Redshift clusters

Query data using Redshift query editor

Use the query editor v2 to run queries in your Redshift cluster.

Query data

Work with your client tools

You can connect to Amazon Redshift from your client tools, such as SQL clients, business intelligence (BI) tools, and extract, transform, load (ETL) tools, using JDBC or ODBC drivers.

Cluster

Cluster identifier

Copy JDBC URL Copy ODBC URL

Choose your JDBC or ODBC driver

Use JDBC or ODBC drivers to connect to Amazon Redshift from your client tools, such as SQL clients, BI tools, and ETL tools. We recommend using the new Amazon Redshift-specific drivers for better performance and scalability.

Driver

JDBC 4.2 without AWS SDK (jar)

Download driver

Clusters (1) **Info**

🔄 Query data ▼ Actions ▼ Create cluster

🔍 Filter clusters by property or value

<input type="checkbox"/>	Cluster	Cluster namespace	Status	Storage capacity us...	CPU utilization	Snapsh...	Notificati...	Tag
<input type="checkbox"/>	redshift-cluster-1 ra3.xplus 1 node 32 TB	1d826245-98e3-41ac-...	Modifying Creating					

← → ↺

us-east-1.console.aws.amazon.com/redshiftv2/home?region=us-east-1#clusters

🔍 📌 ⭐ ⚙️ 📄 👤 ⋮

poll_form.php

Canny edge detecti...

Simple guide on ho...

A Neural Network P...

Learner Lab

Image Feature Extra...

Canny edge detecti...

How to tune hyper...

Machine Learning u...

»

AWS

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

voclabs/user1591029-peeyapak.somvitoon@g.snu.ac.th @ 0748-312...

Amazon Redshift

×

Redshift serverless [New](#)

Provisioned clusters dashboard

▼ Clusters

Reserved nodes

Snapshots

Query editor

Query editor v2 [🔗](#)

Queries and loads

Datashares

▼ Configurations

Workload management

Subnet groups

HSM

Manage Tags

Advisor

AWS Marketplace

Alarms

Events

What's new

Cluster subnet group cluster-subnet-group-2 was create successfully

×

🔔

Successfully created cluster redshift-cluster-1

To manage clusters or connect to the cluster, go to the [redshift-cluster-1](#) details page.

View connection credentials

×

Successfully loaded sample data

To query the sample data, use the Redshift query editor v2.

Query data

×

Amazon Redshift > Clusters

In my account From other accounts

▼ Connect to Redshift clusters

Query data using Redshift query editor

Use the query editor v2 to run queries in your Redshift cluster.

Query data

Work with your client tools

You can connect to Amazon Redshift from your client tools, such as SQL clients, business intelligence (BI) tools, and extract, transform, load (ETL) tools, using JDBC or ODBC drivers.

Cluster

Cluster identifier

Copy JDBC URL

Copy ODBC URL

Choose your JDBC or ODBC driver

Use JDBC or ODBC drivers to connect to Amazon Redshift from your client tools, such as SQL clients, BI tools, and ETL tools. We recommend using the new Amazon Redshift-specific drivers for better performance and scalability.

Driver

JDBC 4.2 without AWS SDK (jar)

Download driver

Clusters (1) Info

🔄 Query data ▼ Actions ▼ Create cluster

🔍 Filter clusters by property or value

< 1 > ⚙️

☐

Cluster

▲

Cluster namespace ▼

Status ▼

Storage capacity us... ▼

CPU utilization ▼

Snapsh... ▼

Notificati... ▼

Tag

☐

redshift-cluster-1
ra3.xplus | 1 node | 32 TB

1d826245-98e3-41ac-...

🔄 Modifying

-

← → ↺

us-east-1.console.aws.amazon.com/redshiftv2/home?region=us-east-1#dashboard

🔍 📌 ⚙️ 🗑️ 📄 📱 🧑

poll_form.php

Canny edge detecti...

Simple guide on ho...

A Neural Network P...

Learner Lab

Image Feature Extra...

Canny edge detecti...

How to tune hyper...

Machine Learning u...

»

aws

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

voclaba/user1591029-peeyapak.sornvittoon@gs.wu.ac.th @ 0748-312...

Amazon Redshift

×

Redshift serverless **New**

Provisioned clusters dashboard

Clusters

Query editor

Query editor v2

Queries and loads

Datashares

Configurations

Workload management

Subnet groups

HSM

Manage Tags

Advisor

AWS Marketplace

Alarms

Events

What's new

Cluster subnet group cluster-subnet-group-2 was create successfully

×

Get consistent, fast query performance for thousands of concurrent queries and users. [View tutorial](#)

×

Amazon Redshift

>

Provisioned clusters dashboard

Try Amazon Redshift Serverless

Purchase reserved nodes

Create cluster

Resources overview

Resource data for US East (N. Virginia) Region.

Total nodes

1

On-demand nodes

1

Reserved nodes

0

Reserved nodes available

0

(0 of 0 used)

Automated snapshots

0

Manual snapshots

0

Cluster overview (1)

Any status

Cluster

Status

redshift-cluster-1

Modifying

View all clusters

Cluster metrics

Any clusters

Last hour

View in CloudWatch

Number of queries

Database connections

Disk space used

CPU utilization

Datashares

Authorize other AWS accounts to access datashares created in this AWS account. Associate or decline datashares from other AWS accounts.

Require authorization

0

Require association

0

Alarms (0)

View in CloudWatch

Alarm name

No ongoing alarms

Events (2)

Last 1 week

Date

Event

a minute ago

10:42 PM

Cluster 'redshift-cluster-1' completed restart at 2022-10-10 15:42 UTC.

a minute ago

Amazon Redshift cluster 'redshift-cluster-1'

Amazon Redshift

Redshift serverless New

Provisioned clusters dashboard

Clusters

Query editor

Query editor v2

Queries and loads

Datashares

Configurations

Workload management

Subnet groups

HSM

Manage Tags

Advisor

AWS Marketplace

Alarms

Events

What's new

Cluster subnet group cluster-subnet-group-2 was create successfully

Amazon Redshift query editor v2 is now available
Query editor v2 provides new features such as multistatement query execution, query parameterization, query versioning, visualizations, and query sharing. [Learn more](#) Go to query editor v2

Amazon Redshift > Clusters > redshift-cluster-1

redshift-cluster-1

Actions Edit Add partner integration Query data

General information

Cluster identifier redshift-cluster-1	Status Modifying	Node type ra3.xplus	Endpoint redshift-cluster-1.cbdkb6riwsow.us-east-1.amazonaws.com
Cluster namespace 1d826245-98e3-41ac-8e62-0b6114bd8036	Date created October 10, 2022, 22:41 (UTC+07:00)	Number of nodes 1	JDBC URL jdbc:redshift://redshift-cluster-1.cbdkb6riwsow.us-east-1.amazonaws.com:5439/
	Storage used -		ODBC URL Driver={Amazon Redshift (x64)}; Server=redshift-cluster-1.cbdkb6riwsow.us-east-1.amazonaws.com

Cluster performance Query monitoring Databases Datashares Schedules Maintenance Properties

Recommendations (0)
To improve performance and decrease operating costs, the Amazon Redshift Advisor provides recommendations.

Alarms (0)
CloudWatch alarms are triggered when a metric threshold is met.

Events (2)
Amazon Redshift tracks events that occur on your cluster.

Cluster metrics (11/11)

Search

Last hour Data for every 5 minutes Average

< 1 2 >

Endpoint

[redshift-cluster-1.cbdkb6riwsow.us-east-1.amazonaws.com](#)

us-east-1.console.aws.amazon.com/redshiftv2/home?region=us-east-1#cluster-details?cluster=redshift-cluster-1&tab=performance

poll_form.php Canny edge detecti... Simple guide on ho... A Neural Network P... Learner Lab / Image Feature Extra... Canny edge detecti... How to tune hyper... Machine Learning u... TensorFlow Lite Co... Convolutional Neur...

amazon Services Search for services, features, blogs, docs, and more [Alt+S]

Amazon Redshift

Redshift serverless New

Provisioned clusters dashboard

Clusters

Query editor

Query editor v2

Queries and loads

Datashares

Configurations

Workload management

Subnet groups

HSM

Manage Tags

Advisor

AWS Marketplace

Alarms

Events

What's new

Amazon Redshift > Clusters > redshift-cluster-1

redshift-cluster-1

Actions Edit Add partner integration Query data

General information

Cluster identifier redshift-cluster-1	Status Modifying	Node type ra3.xplus	Endpoint redshift-cluster-1.cbdkb6riwsow.us-east-1.amazonaws.com
Cluster namespace 1d826245-98e3-41ac-8e62-0b6114bd8036	Date created October 10, 2022, 22:41 (UTC+07:00)	Number of nodes 1	JDBC URL jdbc:redshift://redshift-cluster-1.cbdkb6riwsow.us-east-1.amazonaws.com:5439/
	Storage used -		ODBC URL Driver={Amazon Redshift (x64)}; Server=redshift-cluster-1.cbdkb6riwsow.us-east-1.amazonaws.com

Cluster performance Query monitoring Databases Datashares Schedules Maintenance Properties

Recommendations (0)
To improve performance and decrease operating costs, the Amazon Redshift Advisor provides recommendations.

Alarms (0)
CloudWatch alarms are triggered when a metric threshold is met.

Events (2)
Amazon Redshift tracks events that occur on your cluster.

Cluster metrics (11/11)

Search

Last hour Data for every 5 minutes Average

< 1 2 >

CPU utilization
The percentage of CPU utilization.

Percentage disk space used
The percent of disk space used.

Amazon Redshift query editor v2 is now available

Query editor v2 provides new features such as multistatement query execution, query parameterization, query versioning, visualizations, and query sharing. [Learn more](#)

Go to query editor v2

Amazon Redshift > Clusters > redshift-cluster-1

redshift-cluster-1

Actions Edit Add partner integration Query data

General information

Cluster identifier redshift-cluster-1	Status Available	Node type ra3.xplus	Endpoint redshift-cluster-1.cdkb6riwsow-us-east-1.redshift.amazonaws.com:5439...
Cluster namespace 1d826245-98e3-41ac-8e62-0b6114db8036	Date created October 10, 2022, 22:41 (UTC+07:00)	Number of nodes 1	JDBC URL jdbc:redshift://redshift-cluster-1.cdkb6riwsow-us-east-1.redshift.amazo...
	Storage used 0.01% (0.00 of 32 TB used)		ODBC URL Driver={Amazon Redshift (v64)}; Server=redshift-cluster-1.cdkb6riwsow...

Amazon Redshift

Amazon Redshift query editor v2 is now available

Query editor v2 provides new features such as multistatement query execution, query parameterization, query versioning, visualizations, and query sharing. [Learn more](#)

Go to query editor v2

Redshift serverless New

Provisioned clusters dashboard

Clusters

Query editor

Query editor v2

Queries and loads

Datashares

Configurations

Advisor

AWS Marketplace

Alarms

Events

What's new

Amazon Redshift > Query editor

Editor Query history Saved queries Scheduled queries

Connect to a database to run queries and view results.

Resources

Select database

To view schemas, select a database.

Select schema

To view tables, select a schema.

Filter tables

< 1 >

No resources

No resources to display

Query 1 Query 2 Query 3 Query 4

Status database user Connect to database

2 ParquetFileFormat primary key;

3 RunCommand text;

4 Polaris text;

5 Name text;

6 Dev text;

7 App text;

8 ShipBy text;

9 Product text;

10 Location text;

11 Price text;

12 Cabin text;

13 CargoWeight text;

14

Connect and run Save Schedule Clear

Send feedback

Query results Table details

No resources

Connect to a database to run queries and view results.

Connect to database

Connect to database

Connection

Select a recent database connection or create a new database connection.

☐

Use a recent connection

☒

Create a new connection

Authentication

☒

Temporary credentials

Use the GetClusterCredentials IAM permission and your database user to generate temporary access credentials. [Learn more](#)

☐

AWS Secrets Manager

Use a stored secret to authenticate access. [Learn more](#)

Cluster

redshift-cluster-1 (Available)

Database name

dev

Database user

User name authorized to access your database.

awsuser

Cancel

Connect

Amazon Redshift

Redshift serverless [New](#)

Provisioned clusters dashboard

Clusters

Query editor

Query editor v2

Queries and loads

Datashares

Configurations

Advisor

AWS Marketplace

Alarms

Events

What's new

Amazon Redshift query editor v2 is now available

Query editor v2 provides new features such as multistatement query execution, query parameterization, query versioning, visualizations, and query sharing. [Learn more](#)

Go to query editor v2

Amazon Redshift > Query editor

Editor

Query history

Saved queries

Scheduled queries

Resources

Select database

dev

Select schema

public

Filter tables

category

date

event

listing

sales

users

venue

Query 1

Query 2

Query 3

Query 4

+

Status

Connected

database

dev

user

awsuser

Change connection

2 PassengerId test primary key.

3 Survived test,

4 Polars test,

5 Name test,

6 Sex test,

7 Age test,

8 SibSp test,

9 Parch test,

10 Ticket test,

11 Fare test,

12 Cabin test,

13 Embarked test

14

Run

Save

Schedule

Clear

Query results

Table details

Send feedback

raw.githubusercontent.com/datsciencedojo/datasets/master/titanic.csv

poll_form.php Canny edge detect... Simple guide on ho... A Neural Network P... Learner Lab Image Feature Extra... Canny edge detect... How to

Save as

Organize New folder

Quick access

Desktop Downloads Documents Pictures This PC 20221008

Name Date modified

Yesterday (4)

Basic_CNN_Drawing_Model_HW3_id641991300... 9/10/2565 17:07

Basic_CNN_Drawing_Model_id64199130054.ipynb 9/10/2565 16:14

Example_Basic_CNN (1) 9/10/2565 16:02

data_cat.npy 9/10/2565 14:07

Last week (58)

CSV cave cost

File name: titanic.csv

Save as type: All files

Encoding: UTF-8

Save Cancel

Ln 892, Col 57 100% Windows (CRLF) UTF-8

PassengerId,Survived,Pclass,Name,Sex,Age,SibSp,ParCh
1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 2117
2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs T
3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2
4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)"
5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,
6,0,3,"Moran, Mr. James",male,0,0,330877,8.4583,Q
7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51
8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,34
9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina B
10,1,2,"Nasser, Mrs. Nicholas (Adele Achem)",female,
11,1,3,"Sandstrom, Miss. Marguerite Rut",female,4,1,
12,1,1,"Bonnell, Miss. Elizabeth",female,58,0,0,1137
13,0,3,"Saunderscock, Mr. William Henry",male,20,0,0,
14,0,3,"Andersson, Mr. Anders Johan",male,39,1,5,347
15,0,3,"Vestrom, Miss. Hulda Amanda Adolfina",female
16,1,2,"Hewlett, Mrs. (Mary D Kingcome) ",female,55,
17,0,3,"Rice, Master. Eugene",male,2,4,1,382652,29.1
18,1,2,"Williams, Mr. Charles Eugene",male,0,0,2443
19,0,3,"Vander Planke, Mrs. Julius (Emelia Maria Van
20,1,3,"Sasselman, Mrs. Fatima",female,0,0,2649,7.
21,0,2,"Fynney, Mr. Joseph J",male,35,0,0,239865,26
22,1,2,"Beesley, Mr. Lawrence",male,34,0,0,248698,13
23,1,3,"McGowan, Miss. Anna "Annie"","",female,15,0,0
24,1,1,"Sloper, Mr. William Thompson",male,28,0,0,11
25,0,3,"Palsson, Miss. Torborg Danira",female,8,3,1,
26,1,3,"Asplund, Mrs. Carl Oscar (Selma Augusta Emil
27,0,3,"Emir, Mr. Farred Chehab",male,0,0,2631,7.22
28,0,1,"Fortune, Mr. Charles Alexander",male,19,3,2,
29,1,3,"O'Dwyer, Miss. Ellen "Nellie"","",female,0,0,
30,0,3,"Todoroff, Mr. Lalio",male,0,0,349216,7.8958
31,0,1,"Urchurt, Don. Manuel E",male,40,0,0,PC 176
32,1,1,"Spencer, Mrs. William Augustus (Marie Eugeni
33,1,3,"Glynn, Miss. Mary Agatha",female,0,0,335677
34,0,2,"Wheadon, Mr. Edward H",male,66,0,0,C.A. 245
35,0,1,"Heyer, Mr. Edgar Joseph",male,28,1,0,PC 176
36,0,1,"Holverson, Mr. Alexander Oskar",male,42,1,0,
37,1,3,"Hamee, Mr. Hanna",male,0,0,2677,7.2292,C
38,0,3,"Cann, Mr. Ernest Charles",male,21,0,0,A./S. 212
39,0,3,"Vander Planke, Miss. Augusta Maria",female,18,2,0,345764,18,,S
40,1,3,"Nicola-Yarred, Miss. Jamila",female,14,1,0,2651,11.2417,,C
41,0,3,"Ahlin, Mrs. Johan (Johanna Persdotter Larsson)",female,40,1,0,7546,9.475,,S
42,0,2,"Turpin, Mrs. William John Robert (Dorothy Ann Wonnacott)",female,27,1,0,11668,21,,S
43,0,3,"Kraeff, Mr. Theodor",male,0,0,349253,7.8958,,C
44,1,2,"Laroche, Miss. Simonne Marie Anne Andree",female,3,1,2,SC/Paris 2123,41.5792,,C
45,1,3,"Devaney, Miss. Margaret Delia",female,19,0,0,330958,7.8792,,Q
46,0,3,"Rogers, Mr. William John",male,0,0,S.C./A.4. 23567,8.05,,S
47,0,3,"Lennon, Mr. Denis",male,1,0,370371,15.5,,Q
48,1,3,"O'Driscoll, Miss. Bridget",female,0,0,14311,7.75,,Q
49,0,3,"Samaan, Mr. Youssef",male,2,0,2662,21.6792,,C
50,0,3,"Arnold-Franchi, Mrs. Josef (Josefine Franchi)",female,18,1,0,349237,17.8,,S
51,0,3,"Panula, Master. Juha Niilo",male,7,4,1,3101295,39.6875,,S
52,0,3,"Nosworthy, Mr. Richard Cater",male,21,0,0,A/4. 39886,7.8,,S
53,1,1,"Harper, Mrs. Henry Sleeper (Myna Haxtun)",female,49,1,0,PC 17572,76.7292,D33,C
54,1,2,"Faunthorpe, Mrs. Lizzie (Elizabeth Anne Wilkinson)",female,29,1,0,2926,26,,S
55,0,1,"Ostby, Mr. Engelhart Cornelius",male,65,0,1,113509,61.9792,B30,C
56,1,1,"Woolner, Mr. Hugh",male,0,0,19947,35.5,C52,S
57,1,2,"Rugg, Miss. Emily",female,21,0,0,C.A. 31026,10.5,,S
58,0,3,"Hewlett, Mrs. (Mary D Kingcome) ",female,55,,S

s3.console.aws.amazon.com/s3/buckets?region=us-east-1

poll_form.php Canny edge detect... Simple guide on ho... A Neural Network P... Learner Lab Image Feature Extra... Canny edge detect... How to tune hyper... Machine Learning tu... TensorFlow Lite Co...

Amazon S3

Learn how to effectively use the S3 Storage Classes.

Account snapshot

Last updated: Oct 9, 2022 by Storage Lens. Metrics are generated every 24 hours. Learn more

Total storage 1.7 MB

Object count 225

Avg. object size 7.8 KB

You can enable advanced metrics in the "Default-account-dashboard" configuration.

View Storage Lens dashboard

Buckets (3) info

Buckets are containers for data stored in S3. Learn more

Find buckets by name

Name	AWS Region	Access	Creation date
aws-enn-resources-074831285365-us-east-1	US East (N. Virginia) us-east-1	Objects can be public	September 24, 2022, 17:24:21 (UTC-07:00)
aws-logs-074831285365-us-east-1	US East (N. Virginia) us-east-1	Objects can be public	September 24, 2022, 17:09:55 (UTC-07:00)
tkititanic	US East (N. Virginia) us-east-1	Bucket and objects not public	September 24, 2022, 17:49:12 (UTC-07:00)

AWS: Launching a Redshift Cluster

วิธีเข้าใช้งาน AWS Management Console ผ่าน [AWS Learner Lab](#)

<https://youtu.be/T8noPec75po>

วิธีสร้าง Amazon Redshift cluster ขึ้นมาใช้งาน

<https://youtu.be/BXBb3E95I6Y>

วิธีการทำ Data Ingestion จาก S3 เข้า Amazon Redshift

ข้อมูล Titanic

<https://github.com/datasciencedojo/datasets/blob/master/titanic.csv>

ตัวอย่างโค้ดที่ใช้ในการสร้างตารางใน Redshift

```
CREATE TABLE IF NOT EXISTS titanic (  
  PassengerId text primary key,  
  Survived text,  
  Pclass text,  
  Name text,  
  Sex text,  
  Age text,  
  SibSp text,  
  Parch text,  
  Ticket text,  
  Fare text,  
  Cabin text,  
  Embarked text  
)
```

ตัวอย่างโค้ดที่ใช้ในการโหลดข้อมูลจาก S3 เข้า Redshift

```
COPY titanic FROM 's3://zkan-swu-labs/titanic.csv'  
CREDENTIALS 'aws_iam_role=arn:aws:iam::377290081649:role/LabRole'  
IGNOREHEADER 1 CSV REGION 'us-east-1'
```

ตัวอย่างโค้ดในการทำ Transformation โดยใช้ SQL

```
INSERT INTO  
  new_table (  
    name  
  )  
SELECT  
  Name  
FROM  
  titanic  
WHERE  
  Name NOT IN (SELECT DISTINCT name FROM new_table)
```

Last modified: Saturday, 10 September 2022, 2:21 PM

Upload [Info](#)

Add the files and folders you want to upload to S3. To upload a file larger than 160GB, use the AWS CLI, AWS SDK or Amazon S3 REST API. [Learn more](#)

Drag and drop files and folders you want to upload here, or choose **Add files**, or **Add folders**.

Files and folders (0)

Remove

Add files

Add folder

All files and folders in this table will be uploaded.

Find by name

<

1

>

	Name	Folder	Type	Size
No files or folders				
You have not chosen any files or folders to upload.				

Upload [Info](#)

Add the files and folders you want to upload to S3. To upload a file larger than 160GB, use the AWS CLI, AWS SDK or Amazon S3 REST API. [Learn more](#)

Drag and drop files and folders you want to upload here, or choose **Add files**, or **Add folders**.

Files and folders (1 Total, 59.8 KB)

Remove

Add files

Add folder

All files and folders in this table will be uploaded.

Find by name

< 1 >

<input type="checkbox"/>	Name	Folder	Type	Size
<input type="checkbox"/>	titanic.csv	-	text/csv	59.8 KB

Destination

Destination

s3://kik-my-titanic

► Destination details

Bucket settings that impact new objects stored in the specified destination.

► Permissions

Grant public access and access to other AWS accounts.

► Properties

Specify storage class, encryption settings, tags, and more.

Cancel

Upload

Upload succeeded
View details below.

Upload: status

Close

The information below will no longer be available after you navigate away from this page.

Summary

Destination
s3://kik-my-titanic

Succeeded
1 file, 59.8 KB (100.00%)

Failed
0 files, 0 B (0%)

Files and folders Configuration

Files and folders (1 Total, 59.8 KB)

Find by name

< 1 >

Name	Folder	Type	Size	Status	Error
titanic.csv	-	text/csv	59.8 KB	Succeeded	-

s3.console.aws.amazon.com/s3/buckets/kik-my-titanic?region=us-east-1&tab=objects

Search for services, features, blogs, docs, and more [Alt+S]

Amazon S3

Buckets

Access Points
Object Lambda Access Points
Multi-Region Access Points
Batch Operations
Access analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards
AWS Organizations settings

Feature spotlight

AWS Marketplace for S3

Learn how to effectively use the S3 Storage Classes.

Learn more

Amazon S3 > Buckets > kik-my-titanic

kik-my-titanic

Info

Objects Properties Permissions Metrics Management Access Points

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Copy S3 URI Copy URL Download Open Delete Actions Create folder Upload

Find objects by prefix

< 1 >

	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	titanic.csv	csv	October 10, 2022, 23:06:23 (UTC+07:00)	59.8 KB	Standard

s3.console.aws.amazon.com/s3/object/kik-my-titanic?region=us-east-1&prefix=titanic.csv

Amazon S3

titanic.csv

Copy S3 URI Download Open Object actions

Properties Permissions Versions

Object overview

Owner	awslabs0w4037941t1647900294	S3 URI	s3://kik-my-titanic/titanic.csv
AWS Region	US East (N. Virginia) us-east-1	Amazon Resource Name (ARN)	arn:aws:s3::kik-my-titanic/titanic.csv
Last modified	October 10, 2022, 23:06:23 (UTC+07:00)	Entity tag (Etag)	f247dc127f8cf83460b5cd216d137b51
Size	59.8 KB	Object URL	https://kik-my-titanic.s3.amazonaws.com/titanic.csv
Type	csv		
Key	titanic.csv		

S3 URI

s3://kik-my-titanic/titanic.csv

Object URL

https://kik-my-titanic.s3.amazonaws.com/titanic.csv

Editor

Query history

Saved queries

Scheduled queries



Resources

Info



Select database

Info

To view schemas, select a database.

dev



Select schema

Info

To view tables, select a schema.

public



Filter tables



1



category



date



event



listing



sales



titanic



passengerid

survivedt

pclass

name

sex

age

sibsp

parch

ticket

fare

cabin

embarked



Query 1



Query 2



Query 3



```
1 CREATE TABLE IF NOT EXISTS titanic (  
2   PassengerId text primary key,  
3   Survivedt text,  
4   Pclass text,  
5   Name text,  
6   Sex text,  
7   Age text,  
8   SibSp text,  
9   Parch text,  
10  Ticket text,  
11  Fare text,  
12  Cabin text,  
13  Embarked text  
14 )
```

Run

Save

Schedule

Clear

Query results

Table details

Query



Completed, started on October 10, 2022 at 23:11:46

ELAPSED TIME: 00 m 15 s

us-east-1.console.aws.amazon.com/iamv2/home?region=us-east-1#/home

poll_form.php Canny edge detecti... Simple guide on ho... A Neural Network P... Learner Lab Image Feature Extra... Canny edge detecti... How to tune hyper... Machine Learning u...

aws Services iam

Identity and Access Management (IAM)

Search IAM

Dashboard

▼ Access management

- User groups
- Users
- Roles
- Policies
- Identity providers
- Account settings

▼ Access reports

- Access analyzer
- Archive rules
- Analyzers
- Settings
- Credential report
- Organization activity

Search results for 'iam'

Services (6) See all 6 results ►

Features (17)

Blogs (1,414)

Documentation (116,959)

Knowledge Articles (30)

Events (7)

Marketplace (365)

Services

- IAM** ☆
Manage access to AWS resources
Top features
[Groups](#) [Users](#) [Roles](#) [Policies](#) [Access Analyzer](#)
- IAM Identity Center (successor to AWS Single Sign-On)** ☆
Manage workforce user access to multiple AWS accounts and cloud applications
- Resource Access Manager** ☆
Share AWS resources with other accounts or AWS Organizations
- Amazon VPC IP Address Manager** ☆
Managed IP address management service

Features See all 17 results ►

- Groups**
IAM feature

▼ Access management

User groups

Users

Roles

Policies

Identity providers

Account settings

us-east-1.console.aws.amazon.com/iamv2/home?region=us-east-1#/roles

poll_form.php Canny edge detect... Simple guide on ho... A Neural Network P... Learner Lab Image Feature Extra... Canny edge detect... How to tune hyper... Machine Learning u... TensorFlow Lite Co... Convolutional Neur... python - Cro

aws Services Search for services, features, blogs, docs, and more [Alt+S] Global voclabs/user1591029=peeyapak.somvittoon@g.swu.ac.th @

Identity and Access Management (IAM)

Search IAM

Dashboard

▼ Access management

- User groups
- Users
- Roles**
- Policies
- Identity providers
- Account settings

▼ Access reports

- Access analyzer
- Archive rules

<input type="checkbox"/>	AWSServiceRoleForOrganizations	AWS Service: organizations (Service-Linked Role)	-
<input type="checkbox"/>	AWSServiceRoleForRedshift	AWS Service: redshift (Service-Linked Role)	37 minutes ago
<input type="checkbox"/>	AWSServiceRoleForSupport	AWS Service: support (Service-Linked Role)	-
<input type="checkbox"/>	AWSServiceRoleForTrustedAdvisor	AWS Service: trustedadvisor (Service-Linked Role)	-
<input type="checkbox"/>	EMR_AutoScaling_DefaultRole	AWS Service: elasticmapreduce, and 1 more.	-
<input type="checkbox"/>	EMR_DefaultRole	AWS Service: elasticmapreduce	16 days ago
<input type="checkbox"/>	EMR_EC2_DefaultRole	AWS Service: ec2	16 days ago
<input type="checkbox"/>	LabRole	AWS Service: ssm, and 61 more.	1 hour ago
<input type="checkbox"/>	myRedshiftRole	AWS Service: redshift	-
<input type="checkbox"/>	robomaker_students	AWS Service: robomaker, and 3 more.	-
<input type="checkbox"/>	vocareum	Account: 962395053709	
<input type="checkbox"/>	voclabs	Account: 962395053709	
<input type="checkbox"/>	vocstartsoft	Account: 962395053709	

[IAM](#) > [Roles](#) > [LabRole](#)

LabRole

Delete

Summary

Edit

Creation date September 03, 2022, 17:54 (UTC+07:00)	ARN arn:aws:iam::074831285365:role/LabRole	Instance profile ARN arn:aws:iam::074831285365:instance-profile/LabInstanceProfile
Last activity 1 hour ago	Maximum session duration 1 hour	

[Permissions](#) | [Trust relationships](#) | [Tags \(1\)](#) | [Access Advisor](#) | [Revoke sessions](#)

ARN

[arn:aws:iam::074831285365:role/LabRole](#)

```
COPY titanic FROM 's3://kik-my-titanic/titanic.csv'
CREDENTIALS 'aws_iam_role=arn:aws:iam::074831285365:role/LabRole'
IGNOREHEADER 1 CSV REGION 'us-east-1'
```

Amazon Redshift

Redshift serverless New

Provisioned clusters dashboard

Clusters

Query editor

Query editor v2

Queries and loads

Datashares

Configurations

Advisor

AWS Marketplace

Resources

Select database

dev

Select schema

public

Filter tables

category

date

event

listing

sales

titanic

passengerid

survivedt

pclass

name

Status

Connected

database

dev

user

awsuser

Change connection

Query 1

Query 2

Query 3

Query 4

+

1 COPY titanic FROM 's3://wik-my-titanic/titanic.csv'

2 CREDENTIALS 'aws_iam_role=arn:aws:iam::07483128365:role/LabRole'

3 IGNOREHEADER 1 CSV REGION 'us-east-1'

Run

Save

Schedule

Clear

Send feedback

Query results

Table details

Query 422

Execution

Data

Visualize

Completed, started on October 10, 2022 at 23:19:47

ELAPSED TIME: 00 m 09 s

Amazon Redshift

Redshift serverless New

Provisioned clusters dashboard

Clusters

Query editor

Query editor v2

Queries and loads

Datashares

Configurations

Advisor

AWS Marketplace

Alarms

Events

What's new

Resources

Select database

dev

Select schema

public

Filter tables

category

date

event

listing

sales

titanic

passengerid

survivedt

pclass

name

sex

age

sibsp

parch

ticket

fare

cabin

embarked

users

venue

Editor

Query history

Saved queries

Scheduled queries

Status

Connected

database

dev

user

awsuser

Change connection

Query 1

Query 2

Query 3

Query 4

+

1 select * from titanic

Run

Save

Schedule

Clear

Send feedback

Query results

Table details

Query 456

Execution

Data

Visualize

Completed, started on October 10, 2022 at 23:21:57

ELAPSED TIME: 00 m 02 s

Rows returned (891)

Export

Search rows

passengerid

survivedt

pclass

name

sex

age

sibsp

parch

ticket

fare

cabin

embarked

1

0

3

Braund, Mr. Owen Harris

male

22

1

0

A/5 21171

7.25

S

2

1

1

Cummings, Mrs. John Bradley (Flournoy)

female

38

1

0

PC 17599

71.2833

C85

C

Upload [Info](#)

Add the files and folders you want to upload to S3. To upload a file larger than 160GB, use the AWS CLI, AWS SDK or Amazon S3 REST API. [Learn more](#)

Drag and drop files and folders you want to upload here, or choose **Add files**, or **Add folders**.

Files and folders (1 Total, 70.2 KB)

Remove

Add files

Add folder

All files and folders in this table will be uploaded.

Find by name

< 1 >

<input type="checkbox"/>	Name ▲	Folder ▼	Type ▼	Size ▼
<input type="checkbox"/>	github_events_01.json	-	application/json	70.2 KB

Destination

Destination

s3://kik-my-titanic

► Destination details

Bucket settings that impact new objects stored in the specified destination.

- Permissions

Grant public access and access to other AWS accounts.
- Properties

Specify storage class, encryption settings, tags, and more.







Cancel



Upload

kik-my-titanic [Info](#)[Objects](#) | [Properties](#) | [Permissions](#) | [Metrics](#) | [Management](#) | [Access Points](#)



Objects (2)





Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

  Copy S3 URI  Copy URL  Download  Open  Delete  Actions ▼

 Create folder  Upload

< 1 > 

<input type="checkbox"/>	Name ▲	Type ▼	Last modified ▼	Size ▼	Storage class ▼
<input type="checkbox"/>	 github_events_01.json	json	October 10, 2022, 23:39:10 (UTC+07:00)	70.2 KB	Standard
<input type="checkbox"/>	 titanic.csv	csv	October 10, 2022, 23:06:23 (UTC+07:00)	59.8 KB	Standard

github_events_01.json [Info](#) Copy S3 URI  Download  Open  Object actions ▼[Properties](#) | [Permissions](#) | [Versions](#)

Object overview


Owner
awslabsc0w4037941t1647900294


AWS Region
US East (N. Virginia) us-east-1

Last modified
October 10, 2022, 23:39:10 (UTC+07:00)


Size
70.2 KB


Type
json

Key
 github_events_01.json

S3 URI
 s3://kik-my-titanic/github_events_01.json

Amazon Resource Name (ARN)
 arn:aws:s3:::kik-my-titanic/github_events_01.json

Entity tag (Etag)
 f86b2073d5616ec6fb716cdf907606c0

Object URL
 https://kik-my-titanic.s3.amazonaws.com/github_events_01.json

Upload Info

Add the files and folders you want to upload to S3. To upload a file larger than 160GB, use the AWS CLI, AWS SDK or Amazon S3 REST API. [Learn more](#)

Drag and drop files and folders you want to upload here, or choose **Add files**, or **Add folders**.

Files and folders (1 Total, 133.0 B)
All files and folders in this table will be uploaded.

Remove

Add files

Add folder

🔍

Find by name

< 1 >

<input type="checkbox"/>	Name ▲	Folder ▼	Type ▼	Size ▼
<input type="checkbox"/>	events_json_path.json	-	application/json	133.0 B

Destination

Destination

s3://kik-my-titanic

▶ Destination details

Bucket settings that impact new objects stored in the specified destination.

▶ Permissions

Grant public access and access to other AWS accounts.

▶ Properties

Specify storage class, encryption settings, tags, and more.

Cancel

Upload

Amazon S3 > Buckets > kik-my-titanic

kik-my-titanic [Info](#)

Objects

Properties

Permissions

Metrics

Management

Access Points

Objects (3)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

< 1 > ⚙

<input type="checkbox"/>	Name ▲	Type ▼	Last modified ▼	Size ▼	Storage class ▼
<input type="checkbox"/>	events_json_path.json	json	October 10, 2022, 23:44:02 (UTC+07:00)	133.0 B	Standard
<input type="checkbox"/>	github_events_01.json	json	October 10, 2022, 23:39:10 (UTC+07:00)	70.2 KB	Standard
<input type="checkbox"/>	titanic.csv	csv	October 10, 2022, 23:06:23 (UTC+07:00)	59.8 KB	Standard

Amazon S3 > Buckets > kik-my-titanic > events_json_path.json

events_json_path.json [Info](#)

Properties

Permissions

Versions

Object overview

Owner

awslabsc0w4037941t1647900294

AWS Region

US East (N. Virginia) us-east-1

Last modified

October 10, 2022, 23:44:02 (UTC+07:00)

Size

133.0 B

Type

json

Key

events_json_path.json

S3 URI

s3://kik-my-titanic/events_json_path.json

Amazon Resource Name (ARN)

arn:aws:s3:::kik-my-titanic/events_json_path.json

Entity tag (Etag)

bba012bf51ddc375f18626e7f1e0ce4d

Object URL

https://kik-my-titanic.s3.amazonaws.com/events_json_path.json

```

1 CREATE TABLE IF NOT EXISTS staging_events (
2     id text,
3     type text,
4     actor text,
5     repo text,
6     created_at text
7 )

```

Select database [Info](#)

To view schemas, select a database.

dev ▼

Select schema [Info](#)

To view tables, select a schema.

public ▼

< 1 >

▶ category ...

▶ date ...

▶ event ...

▶ listing ...

▶ sales ...

▶ staging_events ...

▼ titanic ...

passengerid

survivedt

pclass

Query 1 X

Query 2 X

Query 6 X

+

▼

↶ ↷ @ ⚙ /* 📄 🗑

```

1 CREATE TABLE IF NOT EXISTS staging_events (
2     id text,
3     type text,
4     actor text,
5     repo text,
6     created_at text
7 )

```

Run Save Schedule Clear

Send feedback

Query results Table details

Query

Execution

Data

Visualize

Completed, started on October 10, 2022 at 23:47:21

ELAPSED TIME: 00 m 06 s

aws Services iam

Select database [Info](#)
To view schemas, select a database.
dev

Select schema [Info](#)
To view tables, select a schema.
public

Filter tables
< 1 >

- category
- date
- event
- listing
- sales
- staging_events
- titanic
 - passengerid
 - survivedt
 - nrlax

```
1 COPY staging_events FROM 's3://kik-my-titanic/github_events_01.json'
2 CREDENTIALS 'aws_iam_role=arn:aws:iam::074891285365:role/LabRole'
3 JSON 's3://kik-my-titanic/events_json_path.json' REGION 'us-east-1'
4
5
```

Run Save Schedule Clear

Send feedback

Query results Table details

Query 760 [🔗](#)

Execution Data Visualize

Completed, started on October 10, 2022 at 23:48:17
ELAPSED TIME: 00 m 09 s

To confirm deletion, type *permanently delete* in the text input field.

permanently delete

Buckets (3) [Info](#)

Copy ARN

Empty

Delete

Create bucket

Find buckets by name

< 1 > ⚙

	Name ▲	AWS Region ▼	Access ▼	Creation date ▼
<input type="radio"/>	aws-emr-resources-074831285365-us-east-1	US East (N. Virginia) us-east-1	Objects can be public	September 24, 2022, 17:24:21 (UTC+07:00)
<input type="radio"/>	aws-logs-074831285365-us-east-1	US East (N. Virginia) us-east-1	Objects can be public	September 24, 2022, 17:09:35 (UTC+07:00)
<input checked="" type="radio"/>	kiktitanic	US East (N. Virginia) us-east-1	Bucket and objects not public	September 24, 2022, 17:49:12 (UTC+07:00)

Delete redshift-cluster-1?

✕

Deleting the cluster causes the following results:

Deletes all databases (and data) in the cluster.

Deletes the automated snapshot.

Retains all manual snapshots until you manually delete them (none exist).

You can't rotate keys for encrypted manual snapshots if you delete this cluster.

Removes access to the data in datashares for data consumers, including subscribers.

Are you sure that you want to permanently delete **redshift-cluster-1**?

Final snapshot

You can create a final manual snapshot of your cluster before it's deleted so you can later restore it. Restoring it enables you to resume running the cluster and querying data.

☐

Create final snapshot

To confirm deletion, enter *delete* in the field and choose Delete.

delete

Cancel

Delete cluster